

Detection-of-Manipulated-and-Authentic-Images

- There are 4 dataset:
 - Dataset 1: [Link](#)
 - Dataset 2: [Link](#)
 - Dataset 3: [Link](#)
 - Dataset 4: [Link](#)
- The datasets is divided into three sets: training, testing, and validation. Each set contains 'real' and 'fake' images.

Import Libraries

```
import os
import numpy as np
import pandas as pd
import seaborn as sns
import tensorflow as tf
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report, confusion_matrix
```

Global Constants

```
IMG_SIZE = (256, 256)
BATCH_SIZE = 32
AUTOTUNE = tf.data.AUTOTUNE
EPOCHS = 20
```

Load Dataset Function

```
def load_datasets(data_path):
    print(f"\n--- Loading Data from: {data_path} ---")

    # Check if directories exist
    for sub in ['train', 'validation', 'test']:
        if not os.path.exists(os.path.join(data_path, sub)):
            print(f"Error: Directory '{sub}' not found in {data_path}")
            return None, None, None

    train_ds = tf.keras.utils.image_dataset_from_directory(
        os.path.join(data_path, 'train'),
```

```

        labels='inferred',
        label_mode='binary',
        image_size=IMG_SIZE,
        interpolation='nearest',
        batch_size=BATCH_SIZE,
        shuffle=True
    )

val_ds = tf.keras.utils.image_dataset_from_directory(
    os.path.join(data_path, 'validation'),
    labels='inferred',
    label_mode='binary',
    image_size=IMG_SIZE,
    interpolation='nearest',
    batch_size=BATCH_SIZE,
    shuffle=False
)

test_ds = tf.keras.utils.image_dataset_from_directory(
    os.path.join(data_path, 'test'),
    labels='inferred',
    label_mode='binary',
    image_size=IMG_SIZE,
    interpolation='nearest',
    batch_size=BATCH_SIZE,
    shuffle=False
)

# Optimize dataset performance
train_ds = train_ds.cache().prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
test_ds = test_ds.prefetch(buffer_size=AUTOTUNE)

return train_ds, val_ds, test_ds

```

Build Model Function

```

def build_model(input_shape=(256, 256, 3)):
    inputs = tf.keras.Input(shape=input_shape)

    # Data Augmentation & Preprocessing
    x = tf.keras.layers.RandomFlip("horizontal")(inputs)
    x = tf.keras.layers.RandomRotation(0.1)(x)
    x = tf.keras.layers.RandomZoom(0.1)(x)
    x = tf.keras.layers.Rescaling(1./255)(x)

    # CNN Blocks
    # Block 1
    x = tf.keras.layers.Conv2D(32, 3, padding='same',

```

```

activation='relu'))(x)
    x = tf.keras.layers.BatchNormalization()(x)
    x = tf.keras.layers.MaxPooling2D()(x)
    x = tf.keras.layers.Dropout(0.2)(x)

    # Block 2
    x = tf.keras.layers.Conv2D(64, 3, padding='same',
activation='relu'))(x)
    x = tf.keras.layers.BatchNormalization()(x)
    x = tf.keras.layers.MaxPooling2D()(x)
    x = tf.keras.layers.Dropout(0.2)(x)

    # Block 3
    x = tf.keras.layers.Conv2D(128, 3, padding='same',
activation='relu'))(x)
    x = tf.keras.layers.BatchNormalization()(x)
    x = tf.keras.layers.MaxPooling2D()(x)
    x = tf.keras.layers.Dropout(0.3)(x)

    # Block 4
    x = tf.keras.layers.Conv2D(256, 3, padding='same',
activation='relu'))(x)
    x = tf.keras.layers.BatchNormalization()(x)
    x = tf.keras.layers.MaxPooling2D()(x)
    x = tf.keras.layers.Dropout(0.4)(x)

    # Classifier
    x = tf.keras.layers.Flatten()(x)
    x = tf.keras.layers.Dense(512, activation='relu'))(x)
    x = tf.keras.layers.BatchNormalization()(x)
    x = tf.keras.layers.Dropout(0.5)(x)

outputs = tf.keras.layers.Dense(1, activation='sigmoid'))(x)

model = tf.keras.Model(inputs, outputs)

model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy']
)

return model

```

Train Model Function

```

def train_model(model, train_ds, val_ds, epochs=EP0CHS):
    callbacks = [
        tf.keras.callbacks.EarlyStopping(

```

```

        monitor='val_loss', patience=5, restore_best_weights=True
    ),
    tf.keras.callbacks.ReduceLROnPlateau(
        monitor='val_loss', factor=0.2, patience=3, min_lr=1e-6
    )
]

history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs,
    callbacks=callbacks,
    verbose=1
)
return history

```

Plotting Function

```

def plot_history(history, dataset_name="Dataset"):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs_range = range(len(acc))

    plt.figure(figsize=(12, 6))

    # Plot Accuracy
    plt.subplot(1, 2, 1)
    plt.plot(epochs_range, acc, label='Training Accuracy')
    plt.plot(epochs_range, val_acc, label='Validation Accuracy')
    plt.legend(loc='lower right')
    plt.title(f'{dataset_name}: Training and Validation Accuracy')

    # Plot Loss
    plt.subplot(1, 2, 2)
    plt.plot(epochs_range, loss, label='Training Loss')
    plt.plot(epochs_range, val_loss, label='Validation Loss')
    plt.legend(loc='upper right')
    plt.title(f'{dataset_name}: Training and Validation Loss')

    plt.show()

```

Evaluate Model Function

```

def evaluate_model(model, test_ds):
    print("\n--- Evaluating Model ---")

```

```

# 1. Basic Evaluate
loss, accuracy = model.evaluate(test_ds)
print(f'Test Loss: {loss:.4f}')
print(f'Test Accuracy: {accuracy:.4f}')

# 2. Get predictions for detailed metrics
y_pred = []
y_true = []

# Iterate over the test dataset to extract labels and predictions
# Note: test_ds is not shuffled, so order is preserved
for images, labels in test_ds:
    y_true.extend(labels.numpy().flatten())
    preds = model.predict(images, verbose=0)
    y_pred.extend(preds.flatten().round())

y_true = np.array(y_true)
y_pred = np.array(y_pred)

# 3. Classification Report
print('\nClassification Report:')
print(classification_report(y_true, y_pred, target_names=['Real (0)', 'Fake (1)']))

# 4. Confusion Matrix
print('Confusion Matrix:')
cm = confusion_matrix(y_true, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Real', 'Fake'], yticklabels=['Real', 'Fake'])
plt.xlabel('Predicted')
plt.ylabel('True')
plt.title('Confusion Matrix')
plt.show()

```

Function Pipeline

```

def run_experiment(data_path, dataset_name="Dataset"):
    print(f"=====")
    print(f"STARTING EXPERIMENT: {dataset_name}")
    print(f"=====")

    # 1. Load Data
    train_ds, val_ds, test_ds = load_datasets(data_path)

    if train_ds is None:
        print("Skipping due to data loading error.")
        return

```

```

# 2. Build Model
model = build_model(input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3))
model.summary()

# 3. Train
history = train_model(model, train_ds, val_ds)

# 4. Visualize Training
plot_history(history, dataset_name)

# 5. Evaluate
evaluate_model(model, test_ds)

return model, history

```

Define Paths

```

path_dataset_1 = '/kaggle/input/real-and-fake-images-dataset-for-
image-forensics/Data Set 1/Data Set 1'
path_dataset_2 = '/kaggle/input/real-and-fake-images-dataset-for-
image-forensics/Data Set 2/Data Set 2'
path_dataset_3 = '/kaggle/input/real-and-fake-images-dataset-for-
image-forensics/Data Set 3/Data Set 3'
path_dataset_4 = '/kaggle/input/real-and-fake-images-dataset-for-
image-forensics/Data Set 4/Data Set 4'

```

EXECUTION

```

model_1, history_1 = run_experiment(path_dataset_1,
dataset_name="Dataset 1")

=====
STARTING EXPERIMENT: Dataset 1
=====

--- Loading Data from: /kaggle/input/real-and-fake-images-dataset-for-
image-forensics/Data Set 1/Data Set 1 ---
Found 40002 files belonging to 2 classes.

I0000 00:00:1764775034.393256      47 gpu_device.cc:2022] Created
device /job:localhost/replica:0/task:0/device:GPU:0 with 15513 MB
memory: -> device: 0, name: Tesla P100-PCIE-16GB, pci bus id:
0000:00:04.0, compute capability: 6.0

Found 12360 files belonging to 2 classes.
Found 5227 files belonging to 2 classes.

Model: "functional"

```

Layer (type)	Output Shape
Param #	
0 input_layer (InputLayer)	(None, 256, 256, 3)
0 random_flip (RandomFlip)	(None, 256, 256, 3)
0 random_rotation (RandomRotation)	(None, 256, 256, 3)
0 random_zoom (RandomZoom)	(None, 256, 256, 3)
0 rescaling (Rescaling)	(None, 256, 256, 3)
896 conv2d (Conv2D)	(None, 256, 256, 32)
128 batch_normalization (BatchNormalization)	(None, 256, 256, 32)
0 max_pooling2d (MaxPooling2D)	(None, 128, 128, 32)
0 dropout (Dropout)	(None, 128, 128, 32)
18,496 conv2d_1 (Conv2D)	(None, 128, 128, 64)

	batch_normalization_1	(None, 128, 128, 64)	
256	(BatchNormalization)		
0	max_pooling2d_1 (MaxPooling2D)	(None, 64, 64, 64)	
0	dropout_1 (Dropout)	(None, 64, 64, 64)	
73,856	conv2d_2 (Conv2D)	(None, 64, 64, 128)	
512	batch_normalization_2	(None, 64, 64, 128)	
	(BatchNormalization)		
0	max_pooling2d_2 (MaxPooling2D)	(None, 32, 32, 128)	
0	dropout_2 (Dropout)	(None, 32, 32, 128)	
295,168	conv2d_3 (Conv2D)	(None, 32, 32, 256)	
1,024	batch_normalization_3	(None, 32, 32, 256)	
	(BatchNormalization)		
0	max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 256)	
0	dropout_3 (Dropout)	(None, 16, 16, 256)	

	flatten (Flatten)	(None, 65536)	
0			
	dense (Dense)	(None, 512)	
33,554,944			
	batch_normalization_4	(None, 512)	
2,048			
	(BatchNormalization)		
	dropout_4 (Dropout)	(None, 512)	
0			
	dense_1 (Dense)	(None, 1)	
513			

Total params: 33,947,841 (129.50 MB)

Trainable params: 33,945,857 (129.49 MB)

Non-trainable params: 1,984 (7.75 KB)

Epoch 1/20

```
E0000 00:00:1764775065.101023      47 meta_optimizer.cc:966] layout
failed: INVALID_ARGUMENT: Size of values 0 does not match size of
permutation 4 @ fanin shape
inStatefulPartitionedCall/functional_1/dropout_1/stateless_dropout/
SelectV2-2-TransposeNHWCtoNCHW-LayoutOptimizer
I0000 00:00:1764775066.723839      123 cuda_dnn.cc:529] Loaded cuDNN
version 90300
```

```
1251/1251 ━━━━━━━━━━ 177s 135ms/step - accuracy: 0.6167 -
loss: 0.7704 - val_accuracy: 0.5076 - val_loss: 0.9890 -
learning_rate: 0.0010
```

Epoch 2/20

```
1251/1251 ━━━━━━━━━━ 142s 113ms/step - accuracy: 0.8136 -
loss: 0.4027 - val_accuracy: 0.8201 - val_loss: 0.4039 -
learning_rate: 0.0010
```

Epoch 3/20

```
1251/1251 ━━━━━━━━━━ 141s 113ms/step - accuracy: 0.8712 -
loss: 0.2977 - val_accuracy: 0.7901 - val_loss: 0.4504 -
learning_rate: 0.0010
```

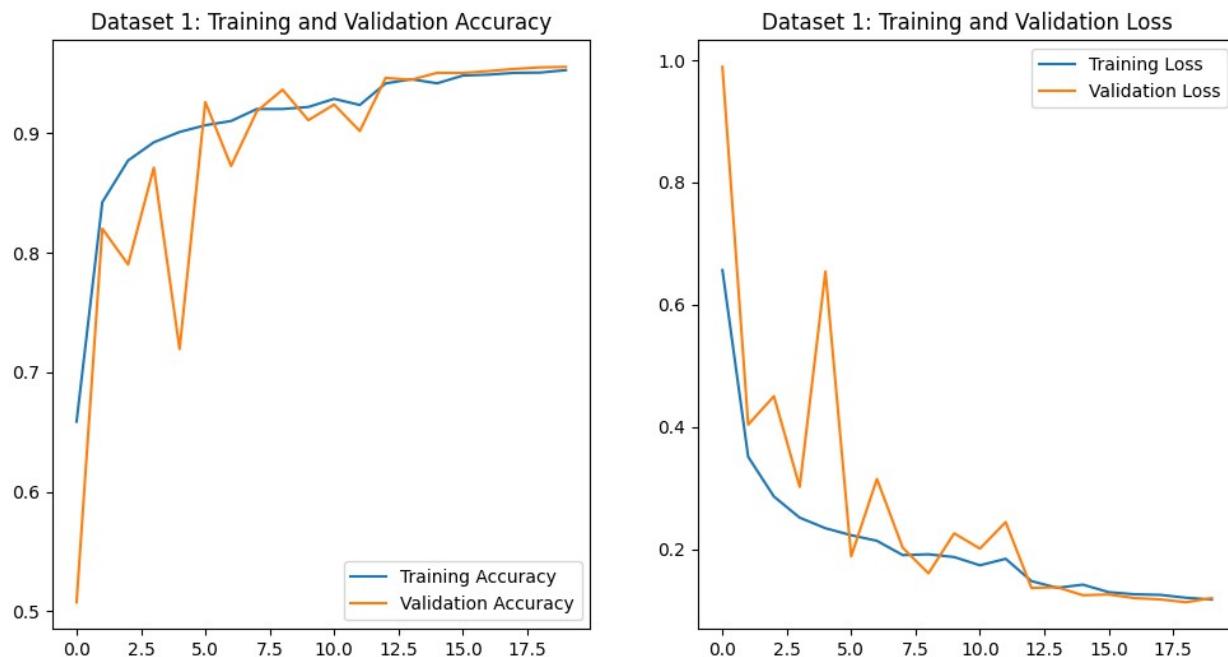
Epoch 4/20

```
1251/1251 ----- 142s 113ms/step - accuracy: 0.8863 -  
loss: 0.2630 - val_accuracy: 0.8712 - val_loss: 0.3021 -  
learning_rate: 0.0010  
Epoch 5/20  
1251/1251 ----- 142s 113ms/step - accuracy: 0.9029 -  
loss: 0.2278 - val_accuracy: 0.7195 - val_loss: 0.6543 -  
learning_rate: 0.0010  
Epoch 6/20  
1251/1251 ----- 141s 113ms/step - accuracy: 0.9047 -  
loss: 0.2268 - val_accuracy: 0.9260 - val_loss: 0.1886 -  
learning_rate: 0.0010  
Epoch 7/20  
1251/1251 ----- 141s 113ms/step - accuracy: 0.9113 -  
loss: 0.2090 - val_accuracy: 0.8726 - val_loss: 0.3150 -  
learning_rate: 0.0010  
Epoch 8/20  
1251/1251 ----- 141s 113ms/step - accuracy: 0.9209 -  
loss: 0.1890 - val_accuracy: 0.9183 - val_loss: 0.2027 -  
learning_rate: 0.0010  
Epoch 9/20  
1251/1251 ----- 141s 113ms/step - accuracy: 0.9197 -  
loss: 0.1920 - val_accuracy: 0.9364 - val_loss: 0.1607 -  
learning_rate: 0.0010  
Epoch 10/20  
1251/1251 ----- 141s 113ms/step - accuracy: 0.9212 -  
loss: 0.1900 - val_accuracy: 0.9108 - val_loss: 0.2261 -  
learning_rate: 0.0010  
Epoch 11/20  
1251/1251 ----- 141s 113ms/step - accuracy: 0.9248 -  
loss: 0.1804 - val_accuracy: 0.9239 - val_loss: 0.2012 -  
learning_rate: 0.0010  
Epoch 12/20  
1251/1251 ----- 141s 113ms/step - accuracy: 0.9246 -  
loss: 0.1811 - val_accuracy: 0.9018 - val_loss: 0.2445 -  
learning_rate: 0.0010  
Epoch 13/20  
1251/1251 ----- 142s 113ms/step - accuracy: 0.9415 -  
loss: 0.1495 - val_accuracy: 0.9462 - val_loss: 0.1368 -  
learning_rate: 2.0000e-04  
Epoch 14/20  
1251/1251 ----- 141s 113ms/step - accuracy: 0.9466 -  
loss: 0.1338 - val_accuracy: 0.9447 - val_loss: 0.1380 -  
learning_rate: 2.0000e-04  
Epoch 15/20  
1251/1251 ----- 142s 113ms/step - accuracy: 0.9386 -  
loss: 0.1502 - val_accuracy: 0.9505 - val_loss: 0.1249 -  
learning_rate: 2.0000e-04  
Epoch 16/20  
1251/1251 ----- 142s 114ms/step - accuracy: 0.9508 -
```

```

loss: 0.1255 - val_accuracy: 0.9503 - val_loss: 0.1261 -
learning_rate: 2.0000e-04
Epoch 17/20
1251/1251 142s 113ms/step - accuracy: 0.9499 -
loss: 0.1271 - val_accuracy: 0.9519 - val_loss: 0.1202 -
learning_rate: 2.0000e-04
Epoch 18/20
1251/1251 142s 113ms/step - accuracy: 0.9508 -
loss: 0.1222 - val_accuracy: 0.9537 - val_loss: 0.1181 -
learning_rate: 2.0000e-04
Epoch 19/20
1251/1251 142s 113ms/step - accuracy: 0.9525 -
loss: 0.1174 - val_accuracy: 0.9550 - val_loss: 0.1134 -
learning_rate: 2.0000e-04
Epoch 20/20
1251/1251 142s 114ms/step - accuracy: 0.9527 -
loss: 0.1161 - val_accuracy: 0.9554 - val_loss: 0.1202 -
learning_rate: 2.0000e-04

```



```

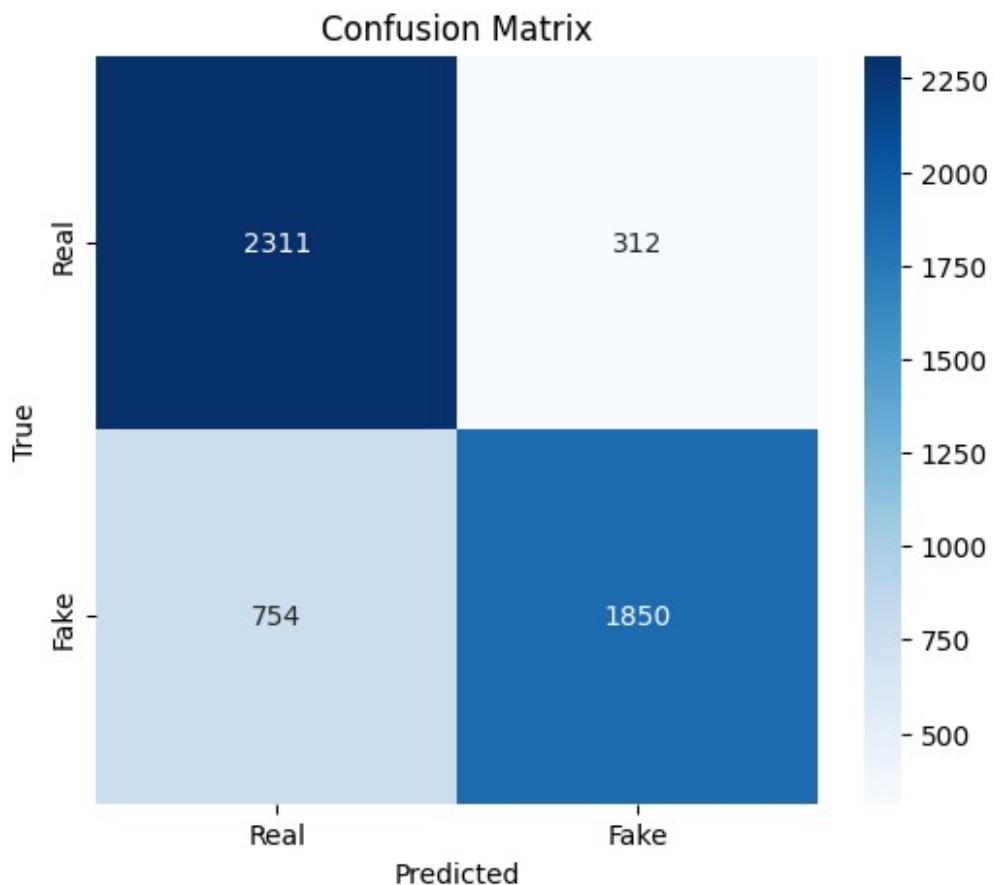
--- Evaluating Model ---
164/164 13s 77ms/step - accuracy: 0.8132 - loss:
0.4866
Test Loss: 0.6067
Test Accuracy: 0.7961

Classification Report:
precision    recall   f1-score   support

```

Real (0)	0.75	0.88	0.81	2623
Fake (1)	0.86	0.71	0.78	2604
accuracy			0.80	5227
macro avg	0.80	0.80	0.79	5227
weighted avg	0.80	0.80	0.79	5227

Confusion Matrix:



```
model_2, history_2 = run_experiment(path_dataset_2,
dataset_name="Dataset 2")
=====
STARTING EXPERIMENT: Dataset 2
=====

--- Loading Data from: /kaggle/input/real-and-fake-images-dataset-for-
image-forensics/Data Set 2/Data Set 2 ---
Found 40000 files belonging to 2 classes.
Found 12356 files belonging to 2 classes.
Found 5226 files belonging to 2 classes.
```

Model: "functional_1"

Layer (type)	Output Shape
Param #	
0 input_layer_1 (InputLayer)	(None, 256, 256, 3)
0 random_flip_1 (RandomFlip)	(None, 256, 256, 3)
0 random_rotation_1 (RandomRotation)	(None, 256, 256, 3)
0 random_zoom_1 (RandomZoom)	(None, 256, 256, 3)
0 rescaling_1 (Rescaling)	(None, 256, 256, 3)
896 conv2d_4 (Conv2D)	(None, 256, 256, 32)
128 batch_normalization_5 (BatchNormalization)	(None, 256, 256, 32)
0 max_pooling2d_4 (MaxPooling2D)	(None, 128, 128, 32)
0 dropout_5 (Dropout)	(None, 128, 128, 32)
18,496 conv2d_5 (Conv2D)	(None, 128, 128, 64)

256	batch_normalization_6 (BatchNormalization)	(None, 128, 128, 64)
0	max_pooling2d_5 (MaxPooling2D)	(None, 64, 64, 64)
0	dropout_6 (Dropout)	(None, 64, 64, 64)
73,856	conv2d_6 (Conv2D)	(None, 64, 64, 128)
512	batch_normalization_7 (BatchNormalization)	(None, 64, 64, 128)
0	max_pooling2d_6 (MaxPooling2D)	(None, 32, 32, 128)
0	dropout_7 (Dropout)	(None, 32, 32, 128)
295,168	conv2d_7 (Conv2D)	(None, 32, 32, 256)
1,024	batch_normalization_8 (BatchNormalization)	(None, 32, 32, 256)
0	max_pooling2d_7 (MaxPooling2D)	(None, 16, 16, 256)
	dropout_8 (Dropout)	(None, 16, 16, 256)

```
0 |  
| flatten_1 (Flatten) | (None, 65536)|  
0 |  
| dense_2 (Dense) | (None, 512)|  
33,554,944 |  
| batch_normalization_9 | (None, 512)|  
2,048 | (BatchNormalization)|  
| dropout_9 (Dropout) | (None, 512)|  
0 |  
| dense_3 (Dense) | (None, 1)|  
513 |
```

Total params: 33,947,841 (129.50 MB)

Trainable params: 33,945,857 (129.49 MB)

Non-trainable params: 1,984 (7.75 KB)

Epoch 1/20

```
E0000 00:00:1764778039.980384      47 meta_optimizer.cc:966] layout  
failed: INVALID_ARGUMENT: Size of values 0 does not match size of  
permutation 4 @ fanin shape  
inStatefulPartitionedCall/functional_1_1/dropout_5_1/stateless_dropout  
/SelectV2-2-TransposeNHWCtoNCHW-LayoutOptimizer
```

```
1250/1250 ━━━━━━━━━━ 171s 134ms/step - accuracy: 0.5665 -  
loss: 0.8294 - val_accuracy: 0.5647 - val_loss: 0.6854 -  
learning_rate: 0.0010
```

Epoch 2/20

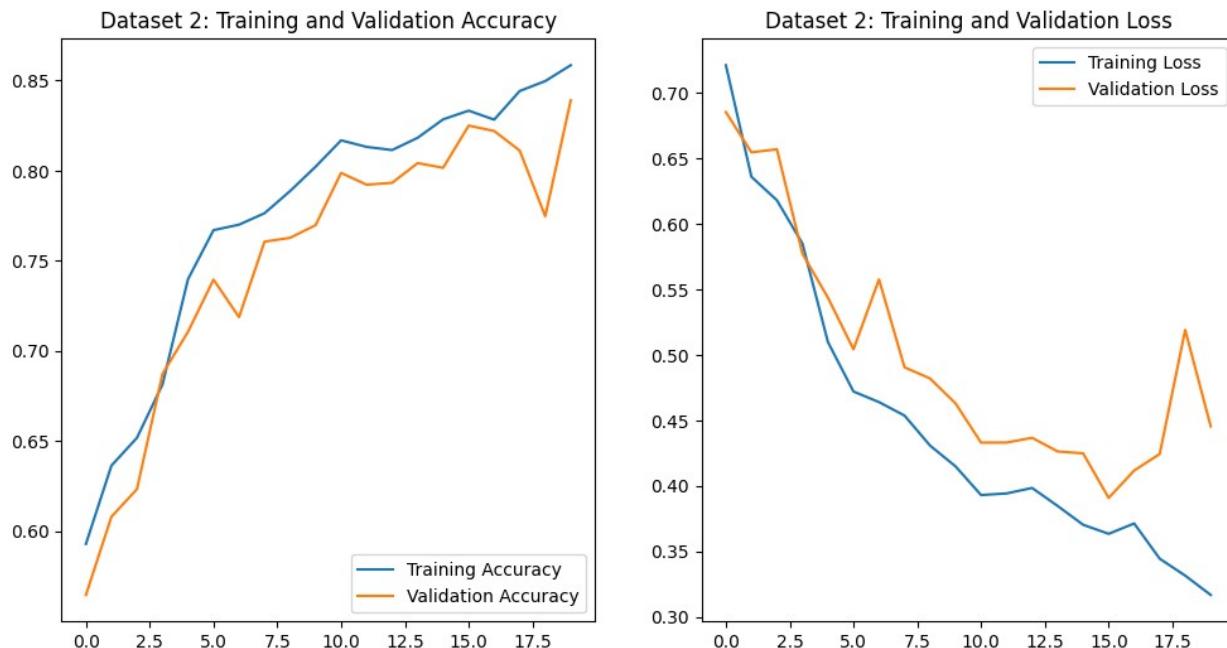
```
1250/1250 ━━━━━━━━━━ 143s 114ms/step - accuracy: 0.6300 -  
loss: 0.6440 - val_accuracy: 0.6081 - val_loss: 0.6547 -  
learning_rate: 0.0010
```

Epoch 3/20

```
1250/1250 ━━━━━━━━━━ 142s 114ms/step - accuracy: 0.6458 -  
loss: 0.6232 - val_accuracy: 0.6234 - val_loss: 0.6571 -  
learning_rate: 0.0010
```

```
Epoch 4/20
1250/1250 ————— 142s 114ms/step - accuracy: 0.6710 -
loss: 0.5997 - val_accuracy: 0.6873 - val_loss: 0.5774 -
learning_rate: 0.0010
Epoch 5/20
1250/1250 ————— 142s 113ms/step - accuracy: 0.7280 -
loss: 0.5257 - val_accuracy: 0.7107 - val_loss: 0.5438 -
learning_rate: 0.0010
Epoch 6/20
1250/1250 ————— 141s 113ms/step - accuracy: 0.7644 -
loss: 0.4732 - val_accuracy: 0.7396 - val_loss: 0.5044 -
learning_rate: 0.0010
Epoch 7/20
1250/1250 ————— 141s 113ms/step - accuracy: 0.7736 -
loss: 0.4566 - val_accuracy: 0.7188 - val_loss: 0.5577 -
learning_rate: 0.0010
Epoch 8/20
1250/1250 ————— 141s 113ms/step - accuracy: 0.7715 -
loss: 0.4617 - val_accuracy: 0.7607 - val_loss: 0.4906 -
learning_rate: 0.0010
Epoch 9/20
1250/1250 ————— 141s 113ms/step - accuracy: 0.7848 -
loss: 0.4318 - val_accuracy: 0.7628 - val_loss: 0.4821 -
learning_rate: 0.0010
Epoch 10/20
1250/1250 ————— 141s 113ms/step - accuracy: 0.7992 -
loss: 0.4188 - val_accuracy: 0.7698 - val_loss: 0.4631 -
learning_rate: 0.0010
Epoch 11/20
1250/1250 ————— 141s 113ms/step - accuracy: 0.8075 -
loss: 0.4069 - val_accuracy: 0.7988 - val_loss: 0.4332 -
learning_rate: 0.0010
Epoch 12/20
1250/1250 ————— 141s 113ms/step - accuracy: 0.8134 -
loss: 0.3926 - val_accuracy: 0.7922 - val_loss: 0.4332 -
learning_rate: 0.0010
Epoch 13/20
1250/1250 ————— 141s 113ms/step - accuracy: 0.8062 -
loss: 0.4071 - val_accuracy: 0.7933 - val_loss: 0.4368 -
learning_rate: 0.0010
Epoch 14/20
1250/1250 ————— 141s 113ms/step - accuracy: 0.8137 -
loss: 0.3925 - val_accuracy: 0.8042 - val_loss: 0.4264 -
learning_rate: 0.0010
Epoch 15/20
1250/1250 ————— 142s 114ms/step - accuracy: 0.8277 -
loss: 0.3704 - val_accuracy: 0.8016 - val_loss: 0.4250 -
learning_rate: 0.0010
Epoch 16/20
```

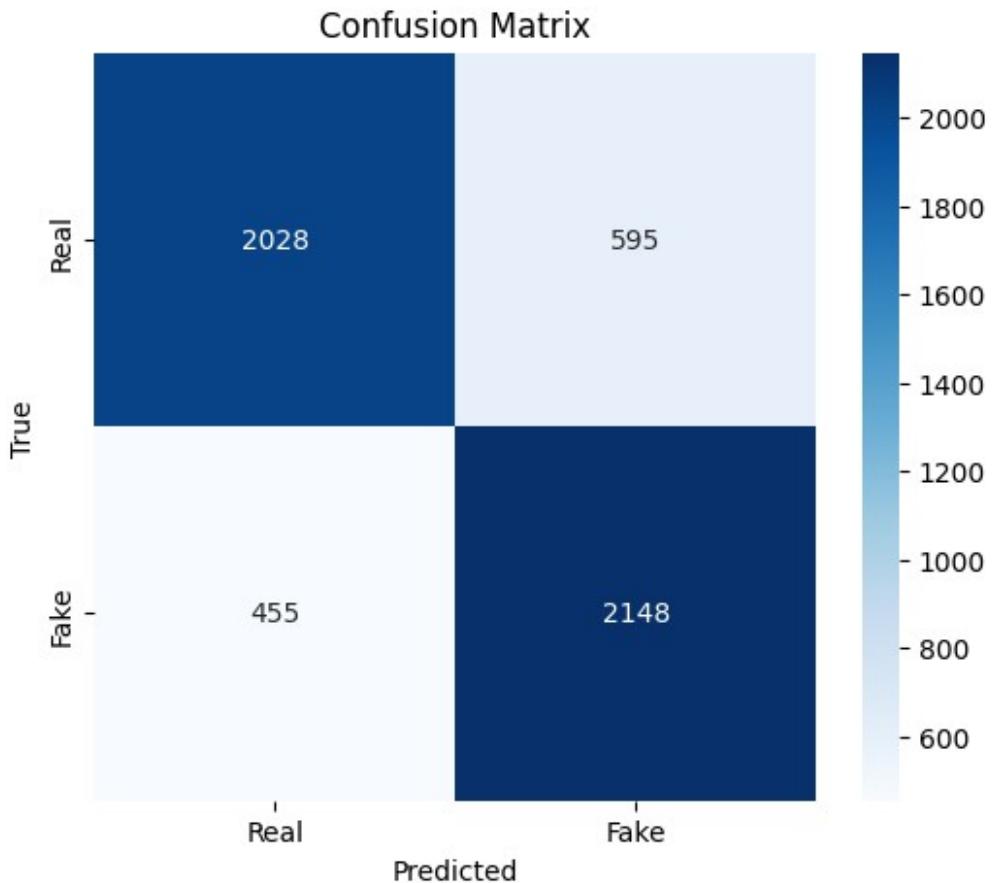
```
1250/1250 ━━━━━━━━━━ 141s 113ms/step - accuracy: 0.8318 -  
loss: 0.3660 - val_accuracy: 0.8250 - val_loss: 0.3909 -  
learning_rate: 0.0010  
Epoch 17/20  
1250/1250 ━━━━━━━━━━ 141s 113ms/step - accuracy: 0.8293 -  
loss: 0.3669 - val_accuracy: 0.8220 - val_loss: 0.4119 -  
learning_rate: 0.0010  
Epoch 18/20  
1250/1250 ━━━━━━━━━━ 141s 113ms/step - accuracy: 0.8404 -  
loss: 0.3492 - val_accuracy: 0.8112 - val_loss: 0.4245 -  
learning_rate: 0.0010  
Epoch 19/20  
1250/1250 ━━━━━━━━━━ 141s 113ms/step - accuracy: 0.8489 -  
loss: 0.3334 - val_accuracy: 0.7748 - val_loss: 0.5192 -  
learning_rate: 0.0010  
Epoch 20/20  
1250/1250 ━━━━━━━━━━ 141s 113ms/step - accuracy: 0.8522 -  
loss: 0.3288 - val_accuracy: 0.8390 - val_loss: 0.4455 -  
learning_rate: 2.0000e-04
```



```
--- Evaluating Model ---  
164/164 ━━━━━━━━━━ 13s 78ms/step - accuracy: 0.7832 - loss:  
0.4429  
Test Loss: 0.4338  
Test Accuracy: 0.7991  
Classification Report:  
precision recall f1-score support
```

Real (0)	0.82	0.77	0.79	2623
Fake (1)	0.78	0.83	0.80	2603
accuracy			0.80	5226
macro avg	0.80	0.80	0.80	5226
weighted avg	0.80	0.80	0.80	5226

Confusion Matrix:



```
model_3, history_3 = run_experiment(path_dataset_3,
dataset_name="Dataset 3")
```

```
=====
STARTING EXPERIMENT: Dataset 3
=====
```

```
--- Loading Data from: /kaggle/input/real-and-fake-images-dataset-for-
image-forensics/Data Set 3/Data Set 3 ---
Found 40000 files belonging to 2 classes.
Found 12356 files belonging to 2 classes.
Found 5226 files belonging to 2 classes.
```

Model: "functional_2"

Layer (type)	Output Shape
Param #	
0 input_layer_2 (InputLayer)	(None, 256, 256, 3)
0 random_flip_2 (RandomFlip)	(None, 256, 256, 3)
0 random_rotation_2 (RandomRotation)	(None, 256, 256, 3)
0 random_zoom_2 (RandomZoom)	(None, 256, 256, 3)
0 rescaling_2 (Rescaling)	(None, 256, 256, 3)
896 conv2d_8 (Conv2D)	(None, 256, 256, 32)
128 batch_normalization_10 (BatchNormalization)	(None, 256, 256, 32)
0 max_pooling2d_8 (MaxPooling2D)	(None, 128, 128, 32)
0 dropout_10 (Dropout)	(None, 128, 128, 32)
18,496 conv2d_9 (Conv2D)	(None, 128, 128, 64)

256	batch_normalization_11 (BatchNormalization)	(None, 128, 128, 64)
0	max_pooling2d_9 (MaxPooling2D)	(None, 64, 64, 64)
0	dropout_11 (Dropout)	(None, 64, 64, 64)
73,856	conv2d_10 (Conv2D)	(None, 64, 64, 128)
512	batch_normalization_12 (BatchNormalization)	(None, 64, 64, 128)
0	max_pooling2d_10 (MaxPooling2D)	(None, 32, 32, 128)
0	dropout_12 (Dropout)	(None, 32, 32, 128)
295,168	conv2d_11 (Conv2D)	(None, 32, 32, 256)
1,024	batch_normalization_13 (BatchNormalization)	(None, 32, 32, 256)
0	max_pooling2d_11 (MaxPooling2D)	(None, 16, 16, 256)
	dropout_13 (Dropout)	(None, 16, 16, 256)

```
0 |  
| flatten_2 (Flatten) | (None, 65536)|  
0 |  
| dense_4 (Dense) | (None, 512)|  
33,554,944 |  
| batch_normalization_14 | (None, 512)|  
2,048 | (BatchNormalization)|  
| dropout_14 (Dropout) | (None, 512)|  
0 |  
| dense_5 (Dense) | (None, 1)|  
513 |
```

Total params: 33,947,841 (129.50 MB)

Trainable params: 33,945,857 (129.49 MB)

Non-trainable params: 1,984 (7.75 KB)

Epoch 1/20

```
E0000 00:00:1764781222.157465      47 meta_optimizer.cc:966] layout  
failed: INVALID_ARGUMENT: Size of values 0 does not match size of  
permutation 4 @ fanin shape  
inStatefulPartitionedCall/functional_2_1/dropout_10_1/stateless_dropout/SelectV2-2-TransposeNHWCToNCHW-LayoutOptimizer
```

```
1250/1250 ━━━━━━━━━━ 167s 130ms/step - accuracy: 0.6391 -  
loss: 0.7653 - val_accuracy: 0.5855 - val_loss: 0.8953 -  
learning_rate: 0.0010
```

Epoch 2/20

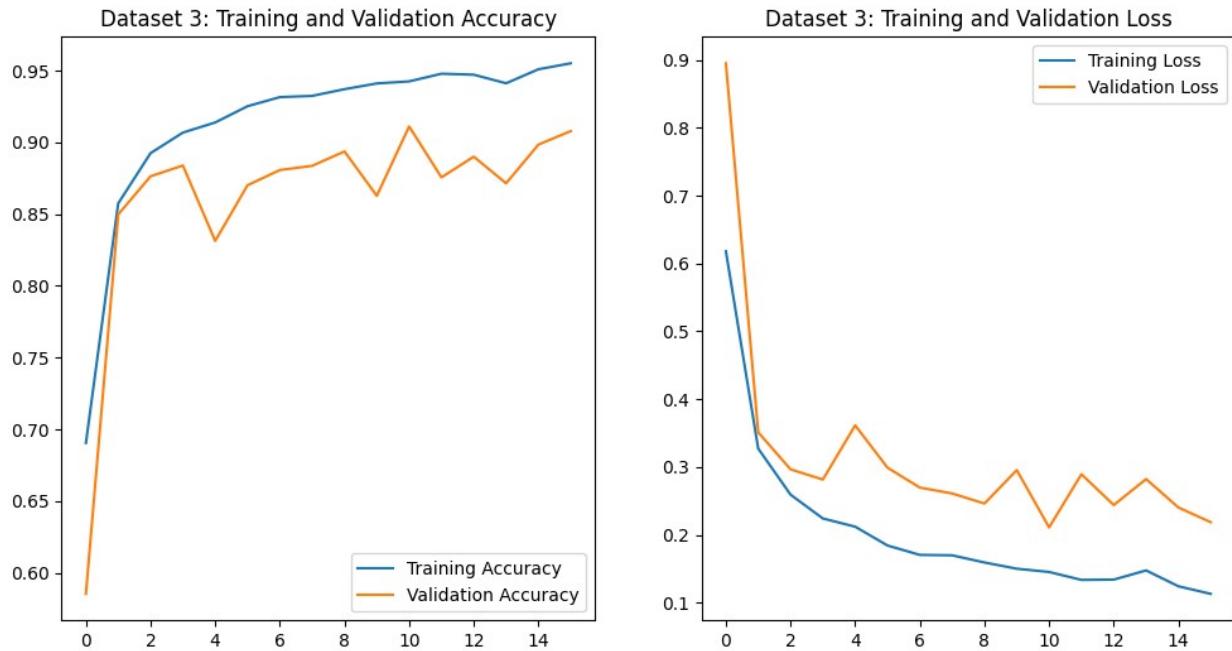
```
1250/1250 ━━━━━━━━━━ 143s 114ms/step - accuracy: 0.8336 -  
loss: 0.3685 - val_accuracy: 0.8498 - val_loss: 0.3511 -  
learning_rate: 0.0010
```

Epoch 3/20

```
1250/1250 ━━━━━━━━━━ 142s 114ms/step - accuracy: 0.8911 -  
loss: 0.2610 - val_accuracy: 0.8763 - val_loss: 0.2966 -  
learning_rate: 0.0010
```

```
Epoch 4/20
1250/1250 ————— 143s 114ms/step - accuracy: 0.9057 -
loss: 0.2282 - val_accuracy: 0.8839 - val_loss: 0.2815 -
learning_rate: 0.0010
Epoch 5/20
1250/1250 ————— 142s 114ms/step - accuracy: 0.9136 -
loss: 0.2150 - val_accuracy: 0.8313 - val_loss: 0.3614 -
learning_rate: 0.0010
Epoch 6/20
1250/1250 ————— 143s 114ms/step - accuracy: 0.9262 -
loss: 0.1825 - val_accuracy: 0.8701 - val_loss: 0.2990 -
learning_rate: 0.0010
Epoch 7/20
1250/1250 ————— 142s 114ms/step - accuracy: 0.9315 -
loss: 0.1695 - val_accuracy: 0.8807 - val_loss: 0.2697 -
learning_rate: 0.0010
Epoch 8/20
1250/1250 ————— 142s 114ms/step - accuracy: 0.9356 -
loss: 0.1614 - val_accuracy: 0.8836 - val_loss: 0.2610 -
learning_rate: 0.0010
Epoch 9/20
1250/1250 ————— 143s 114ms/step - accuracy: 0.9359 -
loss: 0.1626 - val_accuracy: 0.8937 - val_loss: 0.2462 -
learning_rate: 0.0010
Epoch 10/20
1250/1250 ————— 142s 114ms/step - accuracy: 0.9422 -
loss: 0.1454 - val_accuracy: 0.8627 - val_loss: 0.2954 -
learning_rate: 0.0010
Epoch 11/20
1250/1250 ————— 143s 114ms/step - accuracy: 0.9451 -
loss: 0.1430 - val_accuracy: 0.9111 - val_loss: 0.2109 -
learning_rate: 0.0010
Epoch 12/20
1250/1250 ————— 142s 114ms/step - accuracy: 0.9475 -
loss: 0.1355 - val_accuracy: 0.8756 - val_loss: 0.2893 -
learning_rate: 0.0010
Epoch 13/20
1250/1250 ————— 142s 114ms/step - accuracy: 0.9468 -
loss: 0.1342 - val_accuracy: 0.8900 - val_loss: 0.2439 -
learning_rate: 0.0010
Epoch 14/20
1250/1250 ————— 142s 113ms/step - accuracy: 0.9467 -
loss: 0.1345 - val_accuracy: 0.8715 - val_loss: 0.2822 -
learning_rate: 0.0010
Epoch 15/20
1250/1250 ————— 142s 114ms/step - accuracy: 0.9471 -
loss: 0.1319 - val_accuracy: 0.8984 - val_loss: 0.2403 -
learning_rate: 2.0000e-04
Epoch 16/20
```

```
1250/1250 ━━━━━━━━ 142s 114ms/step - accuracy: 0.9545 -  
loss: 0.1148 - val_accuracy: 0.9078 - val_loss: 0.2186 -  
learning_rate: 2.0000e-04
```

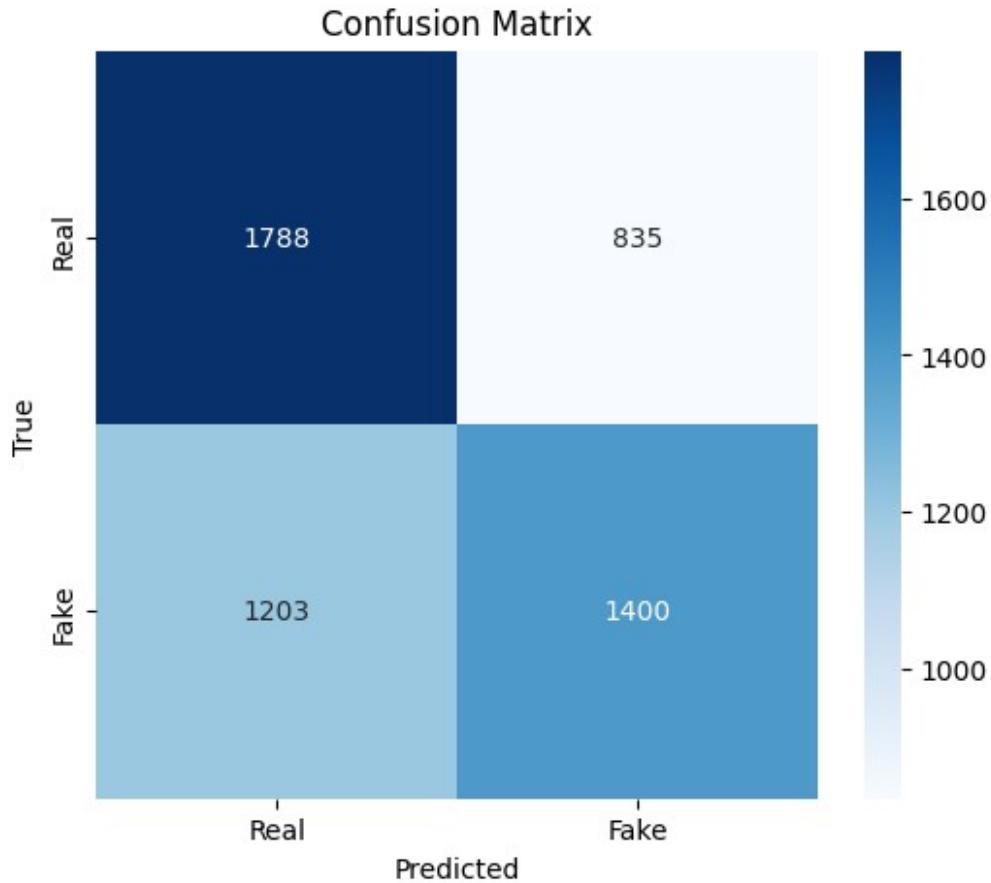


```
--- Evaluating Model ---  
164/164 ━━━━━━━━ 12s 74ms/step - accuracy: 0.5873 - loss:  
1.0412  
Test Loss: 0.9886  
Test Accuracy: 0.6100
```

Classification Report:

	precision	recall	f1-score	support
Real (0)	0.60	0.68	0.64	2623
Fake (1)	0.63	0.54	0.58	2603
accuracy			0.61	5226
macro avg	0.61	0.61	0.61	5226
weighted avg	0.61	0.61	0.61	5226

Confusion Matrix:



```
model_4, history_4 = run_experiment(path_dataset_4,
dataset_name="Dataset 4")
```

```
=====
STARTING EXPERIMENT: Dataset 4
=====
```

```
--- Loading Data from: /kaggle/input/real-and-fake-images-dataset-for-
image-forensics/Data Set 4/Data Set 4 ---
Found 40000 files belonging to 2 classes.
Found 12356 files belonging to 2 classes.
Found 5226 files belonging to 2 classes.
```

```
Model: "functional_4"
```

Layer (type)	Output Shape
Param #	
0 input_layer_4 (InputLayer)	(None, 256, 256, 3)

0	random_flip_4 (RandomFlip)	(None, 256, 256, 3)
0	random_rotation_4 (RandomRotation)	(None, 256, 256, 3)
0	random_zoom_4 (RandomZoom)	(None, 256, 256, 3)
0	rescaling_4 (Rescaling)	(None, 256, 256, 3)
896	conv2d_16 (Conv2D)	(None, 256, 256, 32)
128	batch_normalization_20 (BatchNormalization)	(None, 256, 256, 32)
0	max_pooling2d_16 (MaxPooling2D)	(None, 128, 128, 32)
0	dropout_20 (Dropout)	(None, 128, 128, 32)
18,496	conv2d_17 (Conv2D)	(None, 128, 128, 64)
256	batch_normalization_21 (BatchNormalization)	(None, 128, 128, 64)
	max_pooling2d_17 (MaxPooling2D)	(None, 64, 64, 64)

0		
dropout_21 (Dropout)	(None, 64, 64, 64)	
0		
conv2d_18 (Conv2D)	(None, 64, 64, 128)	
73,856		
batch_normalization_22	(None, 64, 64, 128)	
512	(BatchNormalization)	
max_pooling2d_18 (MaxPooling2D)	(None, 32, 32, 128)	
0		
dropout_22 (Dropout)	(None, 32, 32, 128)	
0		
conv2d_19 (Conv2D)	(None, 32, 32, 256)	
295,168		
batch_normalization_23	(None, 32, 32, 256)	
1,024	(BatchNormalization)	
max_pooling2d_19 (MaxPooling2D)	(None, 16, 16, 256)	
0		
dropout_23 (Dropout)	(None, 16, 16, 256)	
0		
flatten_4 (Flatten)	(None, 65536)	
0		
dense_8 (Dense)	(None, 512)	
33,554,944		

batch_normalization_24 2,048 (BatchNormalization)	(None, 512)	
dropout_24 (Dropout) 0	(None, 512)	
dense_9 (Dense) 513	(None, 1)	

Total params: 33,947,841 (129.50 MB)

Trainable params: 33,945,857 (129.49 MB)

Non-trainable params: 1,984 (7.75 KB)

Epoch 1/20

```
E0000 00:00:1764792379.049701      47 meta_optimizer.cc:966] layout
failed: INVALID_ARGUMENT: Size of values 0 does not match size of
permutation 4 @ fanin shape
inStatefulPartitionedCall/functional_4_1/dropout_20_1/stateless_dropout/SelectV2-2-TransposeNHWCtoNCHW-LayoutOptimizer
```

1250/1250 ██████████ 154s 120ms/step - accuracy: 0.6880 -
loss: 0.6914 - val_accuracy: 0.7883 - val_loss: 0.4502 -
learning_rate: 0.0010

Epoch 2/20

1250/1250 ██████████ 143s 115ms/step - accuracy: 0.8360 -
loss: 0.3636 - val_accuracy: 0.8523 - val_loss: 0.3262 -
learning_rate: 0.0010

Epoch 3/20

1250/1250 ██████████ 143s 114ms/step - accuracy: 0.8975 -
loss: 0.2453 - val_accuracy: 0.8448 - val_loss: 0.3331 -
learning_rate: 0.0010

Epoch 4/20

1250/1250 ██████████ 143s 114ms/step - accuracy: 0.9148 -
loss: 0.2157 - val_accuracy: 0.8788 - val_loss: 0.2711 -
learning_rate: 0.0010

Epoch 5/20

1250/1250 ██████████ 142s 114ms/step - accuracy: 0.9265 -
loss: 0.1831 - val_accuracy: 0.8887 - val_loss: 0.2792 -
learning_rate: 0.0010

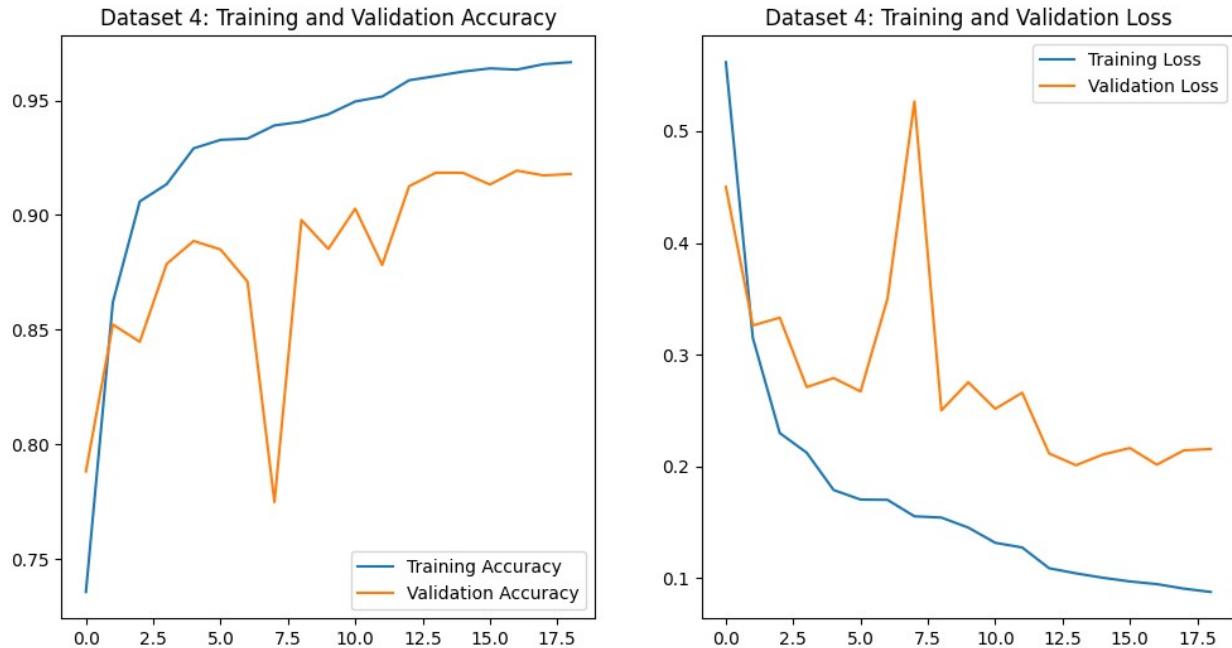
Epoch 6/20

```
1250/1250 ----- 142s 114ms/step - accuracy: 0.9296 -  
loss: 0.1778 - val_accuracy: 0.8850 - val_loss: 0.2671 -  
learning_rate: 0.0010  
Epoch 7/20  
1250/1250 ----- 142s 114ms/step - accuracy: 0.9320 -  
loss: 0.1710 - val_accuracy: 0.8710 - val_loss: 0.3502 -  
learning_rate: 0.0010  
Epoch 8/20  
1250/1250 ----- 142s 114ms/step - accuracy: 0.9342 -  
loss: 0.1639 - val_accuracy: 0.7749 - val_loss: 0.5266 -  
learning_rate: 0.0010  
Epoch 9/20  
1250/1250 ----- 142s 114ms/step - accuracy: 0.9364 -  
loss: 0.1657 - val_accuracy: 0.8979 - val_loss: 0.2502 -  
learning_rate: 0.0010  
Epoch 10/20  
1250/1250 ----- 142s 114ms/step - accuracy: 0.9432 -  
loss: 0.1491 - val_accuracy: 0.8852 - val_loss: 0.2754 -  
learning_rate: 0.0010  
Epoch 11/20  
1250/1250 ----- 142s 114ms/step - accuracy: 0.9470 -  
loss: 0.1355 - val_accuracy: 0.9028 - val_loss: 0.2516 -  
learning_rate: 0.0010  
Epoch 12/20  
1250/1250 ----- 142s 114ms/step - accuracy: 0.9497 -  
loss: 0.1317 - val_accuracy: 0.8783 - val_loss: 0.2660 -  
learning_rate: 0.0010  
Epoch 13/20  
1250/1250 ----- 142s 114ms/step - accuracy: 0.9572 -  
loss: 0.1137 - val_accuracy: 0.9126 - val_loss: 0.2117 -  
learning_rate: 2.0000e-04  
Epoch 14/20  
1250/1250 ----- 142s 114ms/step - accuracy: 0.9586 -  
loss: 0.1072 - val_accuracy: 0.9185 - val_loss: 0.2011 -  
learning_rate: 2.0000e-04  
Epoch 15/20  
1250/1250 ----- 142s 114ms/step - accuracy: 0.9615 -  
loss: 0.1032 - val_accuracy: 0.9184 - val_loss: 0.2108 -  
learning_rate: 2.0000e-04  
Epoch 16/20  
1250/1250 ----- 142s 114ms/step - accuracy: 0.9625 -  
loss: 0.0997 - val_accuracy: 0.9133 - val_loss: 0.2165 -  
learning_rate: 2.0000e-04  
Epoch 17/20  
1250/1250 ----- 142s 114ms/step - accuracy: 0.9623 -  
loss: 0.0973 - val_accuracy: 0.9194 - val_loss: 0.2017 -  
learning_rate: 2.0000e-04  
Epoch 18/20  
1250/1250 ----- 142s 114ms/step - accuracy: 0.9648 -
```

```

loss: 0.0945 - val_accuracy: 0.9173 - val_loss: 0.2144 -
learning_rate: 4.0000e-05
Epoch 19/20
1250/1250 ━━━━━━━━ 142s 114ms/step - accuracy: 0.9642 -
loss: 0.0931 - val_accuracy: 0.9179 - val_loss: 0.2157 -
learning_rate: 4.0000e-05

```



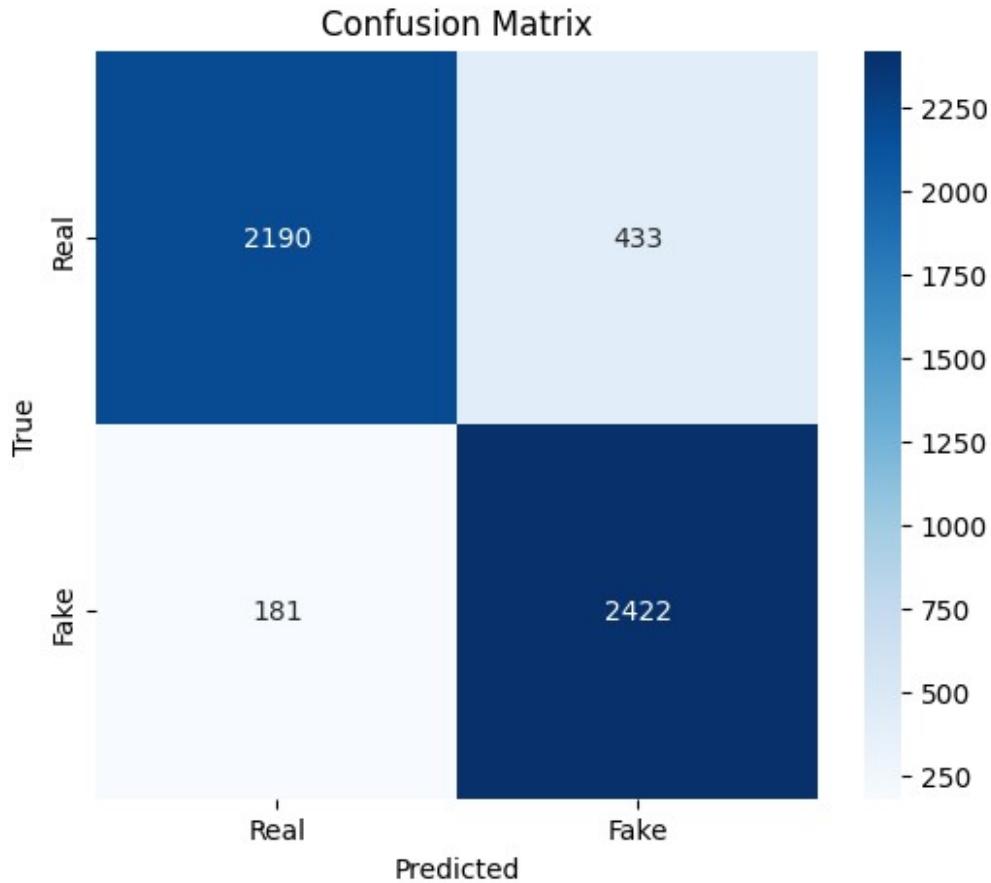
```

--- Evaluating Model ---
164/164 ━━━━━━━━ 5s 33ms/step - accuracy: 0.8532 - loss: 0.3436
Test Loss: 0.2845
Test Accuracy: 0.8825

Classification Report:
precision    recall   f1-score   support
Real (0)      0.92     0.83     0.88     2623
Fake (1)      0.85     0.93     0.89     2603
accuracy          0.88     0.88     0.88     5226
macro avg       0.89     0.88     0.88     5226
weighted avg    0.89     0.88     0.88     5226

Confusion Matrix:

```



Load Combined Datasets

```
def load_combined_datasets(path_list):
    print(f"\n--- Loading and Combining {len(path_list)} Datasets ---")

    full_train_ds = None
    full_val_ds = None
    full_test_ds = None

    for data_path in path_list:
        print(f"Processing: {data_path}...")

        if not os.path.exists(data_path):
            print(f"Warning: Path not found {data_path}, skipping.")
            continue

        # Load partial datasets
        train_ds = tf.keras.utils.image_dataset_from_directory(
            os.path.join(data_path, 'train'),
            labels='inferred',
            label_mode='binary',
```

```

        image_size=IMG_SIZE,
        interpolation='nearest',
        batch_size=BATCH_SIZE,
        shuffle=True
    )

    val_ds = tf.keras.utils.image_dataset_from_directory(
        os.path.join(data_path, 'validation'),
        labels='inferred',
        label_mode='binary',
        image_size=IMG_SIZE,
        interpolation='nearest',
        batch_size=BATCH_SIZE,
        shuffle=False
    )

    test_ds = tf.keras.utils.image_dataset_from_directory(
        os.path.join(data_path, 'test'),
        labels='inferred',
        label_mode='binary',
        image_size=IMG_SIZE,
        interpolation='nearest',
        batch_size=BATCH_SIZE,
        shuffle=False
    )

    # Concatenate logic
    if full_train_ds is None:
        full_train_ds = train_ds
        full_val_ds = val_ds
        full_test_ds = test_ds
    else:
        full_train_ds = full_train_ds.concatenate(train_ds)
        full_val_ds = full_val_ds.concatenate(val_ds)
        full_test_ds = full_test_ds.concatenate(test_ds)

    # 1. Reduce shuffle buffer drastically (from 1000 to 50)
    # This prevents holding too many batches in RAM at once
    full_train_ds = full_train_ds.shuffle(buffer_size=50)

    # 2. REMOVED .cache()
    # We only keep .prefetch(). This loads data from disk as needed.
    full_train_ds = full_train_ds.prefetch(buffer_size=AUTOTUNE)
    full_val_ds = full_val_ds.prefetch(buffer_size=AUTOTUNE)
    full_test_ds = full_test_ds.prefetch(buffer_size=AUTOTUNE)

    print("--- Datasets Combined Successfully ---")
    return full_train_ds, full_val_ds, full_test_ds

```

Function Pipeline

```
def run_experiment(path_list, dataset_name="Combined Datasets"):  
    print(f"=====")  
    print(f"STARTING EXPERIMENT: {dataset_name}")  
    print(f"=====")  
  
    # 1. Load Combined Data  
    train_ds, val_ds, test_ds = load_combined_datasets(path_list)  
  
    if train_ds is None:  
        print("Error: No datasets loaded.")  
        return None, None  
  
    # 2. Build Model  
    model = build_model(input_shape=(IMG_SIZE[0], IMG_SIZE[1], 3))  
  
    # 3. Train  
    history = train_model(model, train_ds, val_ds)  
  
    # 4. Plot  
    plot_history(history, dataset_name)  
  
    # 5. Evaluate  
    evaluate_model(model, test_ds)  
  
    return model, history
```

Define Combined Paths

```
# List of all dataset paths  
all_dataset_paths = [  
  
    '/kaggle/input/real-and-fake-images-dataset-for-image-forensics/Data  
    Set 1/Data Set 1',  
  
    '/kaggle/input/real-and-fake-images-dataset-for-image-forensics/Data  
    Set 2/Data Set 2',  
  
    '/kaggle/input/real-and-fake-images-dataset-for-image-forensics/Data  
    Set 3/Data Set 3',  
  
    '/kaggle/input/real-and-fake-images-dataset-for-image-forensics/Data  
    Set 4/Data Set 4'  
]
```

EXECUTION of Combined Datasets

```
model_combined, history_combined = run_experiment(all_dataset_paths,  
dataset_name="All 4 Datasets Merged")
```

```
=====  
STARTING EXPERIMENT: All 4 Datasets Merged  
=====
```

```
--- Loading and Combining 4 Datasets ---  
Processing: /kaggle/input/real-and-fake-images-dataset-for-image-  
forensics/Data Set 1/Data Set 1...  
Found 40002 files belonging to 2 classes.  
Found 12360 files belonging to 2 classes.  
Found 5227 files belonging to 2 classes.  
Processing: /kaggle/input/real-and-fake-images-dataset-for-image-  
forensics/Data Set 2/Data Set 2...  
Found 40000 files belonging to 2 classes.  
Found 12356 files belonging to 2 classes.  
Found 5226 files belonging to 2 classes.  
Processing: /kaggle/input/real-and-fake-images-dataset-for-image-  
forensics/Data Set 3/Data Set 3...  
Found 40000 files belonging to 2 classes.  
Found 12356 files belonging to 2 classes.  
Found 5226 files belonging to 2 classes.  
Processing: /kaggle/input/real-and-fake-images-dataset-for-image-  
forensics/Data Set 4/Data Set 4...  
Found 40000 files belonging to 2 classes.  
Found 12356 files belonging to 2 classes.  
Found 5226 files belonging to 2 classes.  
--- Datasets Combined Successfully ---
```

```
Epoch 1/20
```

```
E0000 00:00:1764784861.057793      47 meta_optimizer.cc:966] layout  
failed: INVALID_ARGUMENT: Size of values 0 does not match size of  
permutation 4 @ fanin shape  
inStatefulPartitionedCall/functional_3_1/dropout_15_1/stateless_dropout/SelectV2-2-TransposeNHWCToNCHW-LayoutOptimizer
```

```
5001/5001 ━━━━━━━━━━ 623s 123ms/step - accuracy: 0.6941 -  
loss: 0.6049 - val_accuracy: 0.7369 - val_loss: 0.7177 -  
learning_rate: 0.0010
```

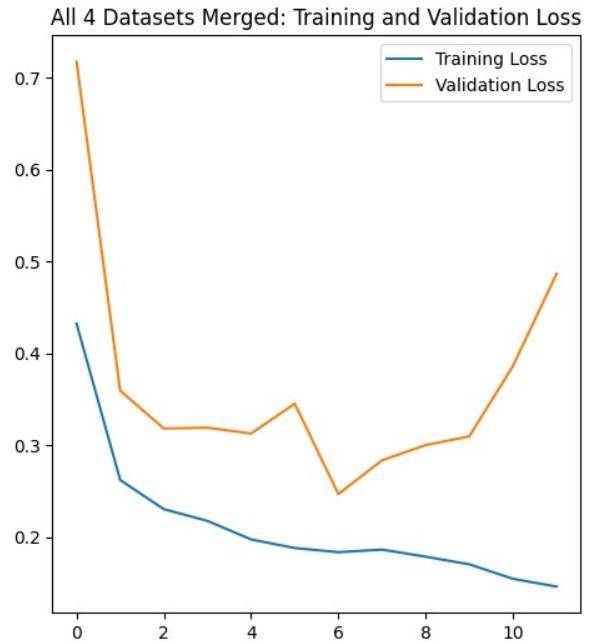
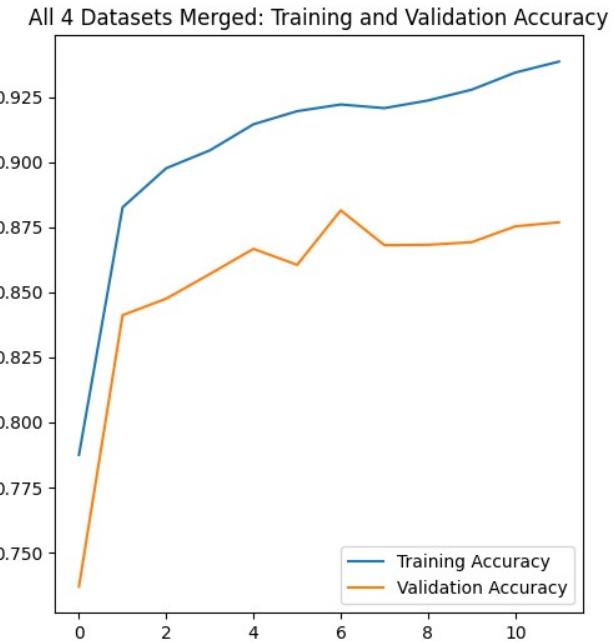
```
Epoch 2/20
```

```
5001/5001 ━━━━━━━━━━ 606s 121ms/step - accuracy: 0.8737 -  
loss: 0.2820 - val_accuracy: 0.8412 - val_loss: 0.3593 -  
learning_rate: 0.0010
```

```
Epoch 3/20
```

```
5001/5001 ━━━━━━━━━━ 599s 119ms/step - accuracy: 0.8931 -  
loss: 0.2412 - val_accuracy: 0.8475 - val_loss: 0.3181 -  
learning_rate: 0.0010
```

```
Epoch 4/20
5001/5001 ━━━━━━━━ 607s 121ms/step - accuracy: 0.8975 -
loss: 0.2313 - val_accuracy: 0.8570 - val_loss: 0.3190 -
learning_rate: 0.0010
Epoch 5/20
5001/5001 ━━━━━━━━ 605s 121ms/step - accuracy: 0.9100 -
loss: 0.2084 - val_accuracy: 0.8667 - val_loss: 0.3126 -
learning_rate: 0.0010
Epoch 6/20
5001/5001 ━━━━━━━━ 608s 121ms/step - accuracy: 0.9153 -
loss: 0.1959 - val_accuracy: 0.8605 - val_loss: 0.3451 -
learning_rate: 0.0010
Epoch 7/20
5001/5001 ━━━━━━━━ 611s 122ms/step - accuracy: 0.9201 -
loss: 0.1877 - val_accuracy: 0.8815 - val_loss: 0.2467 -
learning_rate: 0.0010
Epoch 8/20
5001/5001 ━━━━━━━━ 608s 121ms/step - accuracy: 0.9150 -
loss: 0.1991 - val_accuracy: 0.8681 - val_loss: 0.2834 -
learning_rate: 0.0010
Epoch 9/20
5001/5001 ━━━━━━━━ 609s 121ms/step - accuracy: 0.9173 -
loss: 0.1917 - val_accuracy: 0.8682 - val_loss: 0.3000 -
learning_rate: 0.0010
Epoch 10/20
5001/5001 ━━━━━━━━ 608s 121ms/step - accuracy: 0.9263 -
loss: 0.1735 - val_accuracy: 0.8692 - val_loss: 0.3095 -
learning_rate: 0.0010
Epoch 11/20
5001/5001 ━━━━━━━━ 613s 122ms/step - accuracy: 0.9301 -
loss: 0.1643 - val_accuracy: 0.8753 - val_loss: 0.3860 -
learning_rate: 2.0000e-04
Epoch 12/20
5001/5001 ━━━━━━━━ 608s 121ms/step - accuracy: 0.9345 -
loss: 0.1548 - val_accuracy: 0.8769 - val_loss: 0.4865 -
learning_rate: 2.0000e-04
```



```
--- Evaluating Model ---
656/656 ━━━━━━━━ 28s 43ms/step - accuracy: 0.7988 - loss: 0.4081
Test Loss: 0.4329
Test Accuracy: 0.7835

Classification Report:
precision    recall    f1-score   support
Real (0)      0.76     0.83      0.79    10492
Fake (1)      0.81     0.74      0.77    10413
accuracy          0.79      0.78      0.78    20905
macro avg       0.79      0.78      0.78    20905
weighted avg    0.79      0.78      0.78    20905

Confusion Matrix:
```

Confusion Matrix

